

# BFSI AI Assistant — Technical Documentation

## 1. Introduction

The BFSI AI Assistant is a multi-stage AI system built to handle customer queries in the banking and financial services domain. The primary design goal was to create a solution that balances **accuracy, speed, and reliability** while running fully on local infrastructure.

Instead of relying on a single model for all tasks, my system uses a **tiered architecture** where each stage is optimized for a specific type of query. This approach reduces hallucinations, improves response time, and ensures that policy-related answers are grounded in real documents.

## 2. Design Logic

The system is based on three key principles:

1. Use retrieval whenever possible to ensure correctness
2. Use generation only when reasoning is required
3. Use knowledge grounding for policy or regulatory queries

## 3. Models Used and Rationale

### 3.1 Embedding Model — all-MiniLM-L6-v2

#### Where it is used

- FAQ similarity matching (Tier-1)
- Document retrieval in RAG (Tier-3)

#### Why this model

- Lightweight (fast on CPU)
- Produces high-quality semantic embeddings
- Low memory footprint (ideal for local setup)

#### Role in pipeline

Converts text queries and documents into vector representations so semantic similarity can be computed.

## 3.2 Small Language Model — TinyLlama-1.1B (GGUF Quantized)

### Where it is used

- Tier-2 reasoning
- Tier-3 answer generation

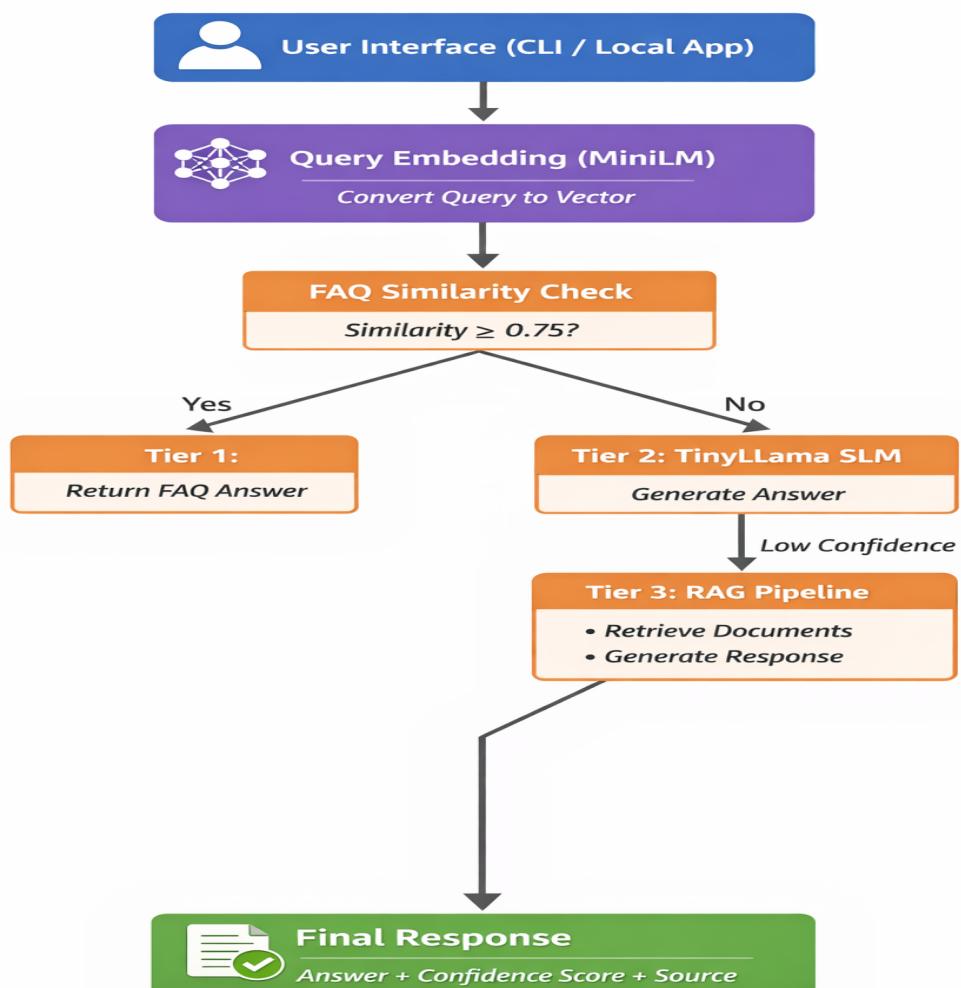
### Why this model

- Small enough to run locally ( $\approx 600$  MB quantized)
- Instruction-tuned for conversational responses
- Low latency compared to larger LLMs

### Role in pipeline

Generates natural language responses when retrieval alone is insufficient.

## 4. Architecture Overview



# 5. Tiered Processing Logic

## 5.1 Tier 1 — FAQ Semantic Retrieval

### Purpose

Provide instant answers for frequently asked questions.

### Process

1. Query converted to embedding
2. Cosine similarity computed with FAQ dataset
3. Highest scoring match selected

### Threshold

SIM\_THRESHOLD = 0.75

### Why 0.75

Empirically balances precision and recall.

Below this value, semantic matches may be too weak to trust.

### Output

Direct deterministic answer from dataset.

## 5.2 Tier 2 — Local Reasoning (SLM)

### Purpose

Handle queries that require procedural explanation but do not need policy grounding.

### Activation Condition

SLM\_THRESHOLD = 0.45

This means:

- Query is somewhat related to known domain
- But not close enough for FAQ retrieval

### Why 0.45

Ensures the model is invoked only when semantic confidence exists, preventing irrelevant generation.

## Prompt Design

The prompt includes:

- System instructions
- User query
- Response constraints

This helps control verbosity and avoid hallucinations.

## 5.3 Tier 3 — Retrieval Augmented Generation

### Purpose

Answer policy, regulatory, and domain-specific queries using knowledge grounding.

### Process

1. Query embedding generated
2. FAISS retrieves top-k documents( $k=3$ )
3. Retrieved context injected into prompt
4. TinyLlama generates a grounded answer

### Why $\text{top\_k} = 3$

Provides enough context without overwhelming the model or increasing latency.

## 6. Vector Database Configuration

**Engine:** FAISS

**Index Type:** Flat L2

### Reason

- Fast similarity search
- Suitable for small-to-medium datasets
- No training required

## 7. Routing Algorithm

The routing decision follows a cascading logic:

1. If FAQ similarity  $\geq 0.75 \rightarrow$  Tier-1
2. Else if similarity  $\geq 0.45 \rightarrow$  Tier-2

3. Else → Tier-3

This ensures:

- Fastest path for simple queries
- Controlled generation for complex queries
- Knowledge grounding for policy queries

## 8. Parameters Summary

Parameter	Value	Purpose
FAQ threshold	0.75	Ensure strong semantic match
SLM threshold	0.45	Trigger reasoning layer
Top-k retrieval	3	Context size for RAG
Embedding dimension	384	Vector representation size
Context window	2048 tokens	LLM input limit

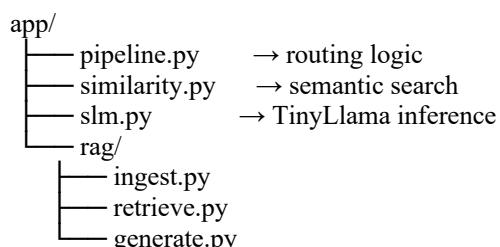
## 9. Performance Profile

The system demonstrates:

- Sub-second responses for FAQ queries
- Moderate latency for LLM responses
- Low memory usage due to quantized model

This makes it suitable for edge or on-prem deployment.

## 10. Implementation Structure



Each module is isolated to maintain separation of concerns.

## 11. Why This Architecture Works Well

- Reduces unnecessary LLM calls
- Improves factual correctness
- Maintains low compute cost

- Provides explainable routing decisions

This is aligned with best practices in enterprise GenAI systems.

## 12. Conclusion

The BFSI AI Assistant demonstrates a practical and scalable approach to building domain-specific AI systems. By combining semantic retrieval, lightweight local language models, and retrieval-augmented generation, the system ensures both efficiency and reliability — two critical requirements in financial applications.